

## Naive Bayes Algorithm

The Naive Bayes Algorithm is a simple method used for classification tasks in machine learning. It works on the principle of Bayes' Theorem, which calculates the probability of a category based on prior knowledge. The algorithm is especially popular for text classification problems such as spam detection, sentiment analysis, and document categorization. It works by determining the probability of each class given a set of features and then predicting the class with the highest probability. By calculating the likelihood of each class and picking the one with the highest probability, the Naive Bayes Algorithm delivers reliable and accurate predictions. Its design, ease of use, and ability to manage both simple and complex classification tasks make it a valuable tool for both beginners and experienced data scientists.

For implementing naive bayes algorithm, we firstly import libraries, and we perform algorithm for IMDB Data dataset.

```
In [1]: import pandas as pd
import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.pipeline import Pipeline
```

```
In [10]: nltk.download('punkt')
nltk.download('stopwords')
nltk.download('wordnet')
import nltk
nltk.download('omw-1.4')

[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\Rubina\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\Rubina\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\Rubina\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package omw-1.4 to
[nltk_data] C:\Users\Rubina\AppData\Roaming\nltk_data...
```

```
Out[10]: True
```

```
In [12]: data = pd.read_csv('IMDB Dataset.csv', encoding='utf-8')
```

This code snippet downloads several natural language processing (NLP) resources using the Natural Language Toolkit (nltk) library in Python. After downloading these resources, the code reads a CSV file named "IMDB Dataset.csv" that is encoded in UTF-8 format using the pandas library (pd). This dataset likely contains movie reviews or related data from IMDB.



After training the sentiment analysis pipeline on IMDB movie reviews, this code snippet evaluates its performance on a test subset ( $X_{\text{test}}$ ,  $y_{\text{test}}$ ). First, it predicts sentiment labels (positive or negative) for the test reviews using `pipeline.predict(X_test)`. The `accuracy_score(y_test, y_pred)` function computes and prints the accuracy of these predictions compared to the actual sentiment labels ( $y_{\text{test}}$ ), indicating how often the model's predictions match the true labels.

Next, the code generates a confusion matrix using `confusion_matrix(y_test, y_pred)`. This matrix provides a detailed breakdown of the model's predictions versus the actual labels, showing the number of true positives, true negatives, false positives, and false negatives. It helps assess the model's performance in correctly classifying positive and negative reviews.

Lastly, the `classification_report(y_test, y_pred)` function generates a report that includes precision, recall, and F1-score metrics for each sentiment class (positive and negative). Precision measures the accuracy of positive predictions, recall indicates the proportion of actual positives correctly identified, and the F1-score is the harmonic mean of precision and recall, offering a balanced assessment of the model's performance for each sentiment category. Together, these outputs provide comprehensive insights into how well the sentiment analysis model performs on unseen data, guiding further refinement or deployment decisions.

```
In [21]: y_pred = pipeline.predict(X_test)
```

```
In [22]: print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

```
Accuracy: 0.8564
```

```
Confusion Matrix:
[[4188  773]
 [ 663 4376]]
```

```
Classification Report:
              precision    recall  f1-score   support

   negative      0.86      0.84      0.85     4961
   positive      0.85      0.87      0.86     5039

   accuracy              0.86              0.86     10000
  macro avg              0.86              0.86     10000
 weighted avg              0.86              0.86     10000
```

After generating confusion matrix, this code predicts whether three new movie reviews are positive or negative using a model trained on IMDB reviews. It first prepares the reviews by making them lowercase, removing unnecessary characters, and transforming them into a format the model can understand. Then, it uses the model to make predictions about the sentiment of each review. Finally, it displays each review alongside its predicted sentiment. This shows how the model can automatically analyze new reviews based on what it learned from past data, providing insights into whether the reviews are positive or negative.

```
In [23]: new_reviews = [  
    "This movie was excellent! I loved it.",  
    "Terrible film. I hated every minute of it.",  
    "An average movie, nothing special."  
]
```

```
In [24]: processed_new_reviews = [preprocess_text(review) for review in new_reviews]  
  
new_predictions = pipeline.predict(processed_new_reviews)  
  
for review, sentiment in zip(new_reviews, new_predictions):  
    print(f"Review: {review}")  
    print(f"Predicted Sentiment: {sentiment}\n")
```

```
Review: This movie was excellent! I loved it.  
Predicted Sentiment: positive
```

```
Review: Terrible film. I hated every minute of it.  
Predicted Sentiment: negative
```

```
Review: An average movie, nothing special.  
Predicted Sentiment: negative
```

## Conclusion

Therefore, we developed and applied a sentiment analysis system for IMDB movie reviews. Beginning with text preprocessing to clean and standardize the data, we trained a model using TF-IDF vectorization and a Multinomial Naive Bayes classifier to predict whether reviews were positive or negative. Evaluating the model's performance involved assessing accuracy, constructing a confusion matrix to visualize errors, and generating a classification report for detailed metrics. Finally, we demonstrated the model's ability to predict sentiments for new reviews, highlighting its effectiveness in automating sentiment analysis tasks and extracting insights from textual data. It helped in analyzing and understanding sentiment in text-based content.